

CSC420 Assignment 1

Gabriel Luong
996268275

1

Algorithm 1 Convolution

```
1: procedure CONVOLUTION(Image( $i \times j$ ), Filter( $u \times v$ ))
2:    $G \leftarrow$  new  $i \times j$  image matrix
3:   for ImageY = 0 to  $j$  do
4:     for ImageX = 0 to  $i$  do
5:       sum  $\leftarrow$  0
6:       for FilterY =  $-v$  to  $v$  do
7:         for FilterX =  $-u$  to  $u$  do
8:           sum  $\leftarrow$  sum + Filter(FilterX, FilterY)  $\times$  Image(ImageX + FilterX, ImageY
+ FilterY)
9:         end for
10:       end for
11:        $G(\text{ImageX}, \text{ImageY}) = \text{sum}$ 
12:     end for
13:   end for
14:   return  $G$ 
15: end procedure
```

2

(a) For a given $n \times n$ image, I , and $m \times m$ filter, h , the process of performing a convolution requires m^2 operations per pixel in I which results in a time complexity of $O(n^2m^2)$. However, the complexity of convolution with the help of Fourier Transform can be reduced to $O(n \log n)$. If h can be separated to vertical and horizontal 1D filters, only 2 1D convolutions (m operations) will need to be computed for every pixel in I , and therefore it will only have a time complexity of $O(2m * n^2) = O(m * n^2)$.

(b) Applying first a Gaussian filter with $\sigma_1 = 3$, and then another Gaussian with $\sigma_2 = 4$ is equivalent to applying one Gaussian filter with $\sigma = \sqrt{\sigma_1^2 + \sigma_2^2} = \sqrt{3^2 + 4^2} = \sqrt{25} = 5$.

(c) Convolve **lena.png** with a Gaussian filter ($\sigma = 1$).

```
1 im = imread('lena.png');
2
3 % create the filter
4 sigma = 1;
5 filter = fspecial('gaussian', sigma * 3, sigma);
6
7 tic;
8 % convolve the image with the Gaussian filter
9 out = imfilter(im, filter, 'same', 'conv');
10 time_conv = toc;
11
12 imshow(out); % show the result of the convolution
13 fprintf('Time for sliding convolution: %0.4f\n', time_conv);
14 % Time for sliding convolution: 0.0020
```



(d) Since Gaussian filter F is separable, you can compute the horizontal and vertical filters, and perform a 1D convolution with the horizontal filter and image to each row followed by a 1D convolution with the vertical filter on each column. This is possible because of the associative property of convolution. You can verify the computed filters by finding the product of the horizontal and vertical filters is equal to the original Gaussian filter ($F = vh^T$). Separability allows for a faster


filtering process since only 2 1D convolutions will need to be computed for every pixel in the image. This is illustrated in the following code with a faster convolution time of 0.001 s.

```
1 im = imread('lena.png');
2
3 % create the Gaussian filter
4 sigma = 1;
5 filter = fspecial('gaussian', sigma * 3, sigma);
6
7 % compute the vertical and horizontal filters of the Gaussian filter
8 [U,S,V] = svd(filter);
9 vertical_filter = sqrt(S(1,1)) * U(:,1);
10 horizontal_filter = sqrt(S(1,1)) * V(:,1)';
11
12 tic;
13 out = conv2(vertical_filter, horizontal_filter, im, 'same');
14 time_conv = toc;
15
16 imshow(out); % show the result of the convolution
17 fprintf('Time for convolution: %0.4f\n', time_conv);
18 % Time for convolution: 0.0010
```

3


(a) Compute the magnitude of gradients for **waldo.png** and **template.png**.

```
1 template = imread('template.png');
2 waldo = imread('waldo.png');
3
4 % compute the magnitude of gradients for the images
5 [template_grad_mag, template_grad_dir] = imgradient(rgb2gray(template));
6 [waldo_grad_mag, waldo_grad_dir] = imgradient(rgb2gray(waldo));
7
8 template_grad_mag = uint8(template_grad_mag);
9 waldo_grad_mag = uint8(waldo_grad_mag);
10
11 figure, imshow(template_grad_mag);
12 figure, imshow(waldo_grad_mag);
```



(b) Localize the template in the image **waldo.png** based on the magnitude of gradients.

```
1 template = imread('template.png');
2 waldo = imread('waldo.png');
3
4 % compute the magnitude of gradients for the images
5 [template-grad-mag, template-grad-dir] = imgradient(rgb2gray(template));
6 [waldo-grad-mag, waldo-grad-dir] = imgradient(rgb2gray(waldo));
7
8 template-grad-mag = uint8(template-grad-mag);
9 waldo-grad-mag = uint8(waldo-grad-mag);
10
11 % Following code reused from lecture 2:
12 % normalized cross-correlation
13 out = normxcorr2(template-grad-mag, waldo-grad-mag);
14
15 % plot the cross-correlation results
16 figure('position', [100,100,size(out,2),size(out,1)]);
17 subplot('position', [0,0,1,1]);
18 imagesc(out)
19 axis off;
20 axis equal;
21
22 % find the peak in response
23 [y,x] = find(out == max(out(:)));
24 y = y(1) - size(template-grad-mag, 1) + 1;
25 x = x(1) - size(template-grad-mag, 2) + 1;
26
27 % plot the detection's bounding box
28 figure('position', [300,100,size(im,2),size(im,1)]);
29 subplot('position', [0,0,1,1]);
30 imshow(waldo)
31 axis off;
32 axis equal;
33 rectangle('position', [x,y,size(template-grad-mag,2),size(template-grad-mag,1)], 'edgecolor',
34         [0.1,0.2,1], 'linewidth', 3.5);
```





(c) Compute the image pyramid with 3 levels for **waldo.png**.

```

1 function out = impyramid(im)
2 % create the Gaussian filter
3 sigma = 1;
4 filter = fspecial('gaussian', sigma * 3, sigma);
5
6 % blur the image with the Gaussian filter
7 blur_im = imfilter(im, filter, 'same', 'conv');
8
9 % separate the channels of an RGB image
10 im_red = blur_im(:,:,1);
11 im_green = blur_im(:,:,2);
12 im_blue = blur_im(:,:,3);
13
14 % downsample the channels by a factor of 2
15 downsample_im_red = im_red(1:2:size(im_red,1), 1:2:size(im_red, 2));
16 downsample_im_green = im_green(1:2:size(im_green,1), 1:2:size(im_green, 2));
17 downsample_im_blue = im_blue(1:2:size(im_blue,1), 1:2:size(im_blue, 2));
18
19 % put the channels back together
20 out = cat(3, downsample_im_red, downsample_im_green, downsample_im_blue);
21 end
22
23 % compute the image pyramid with 3 levels for waldo.png
24 im = imread('waldo.png');
25 im1 = impyramid(im);
26 im2 = impyramid(im1);
27 im3 = impyramid(im2);

```

28

```
29 imshow(im);  
30 figure, imshow(im1);  
31 figure, imshow(im2);  
32 figure, imshow(im3);
```





4

To reduce the amount of fine, detailed edges that are detected with the Canny edge detector, one can (1) raise the threshold of the low threshold, and (2) use a Gaussian filter with a larger σ .

5

(a) Compute the magnitude of gradients for an image.

```
1 im = imread('temple.jpg');
2
3 % compute the magnitude of gradients of the image
4 [imMag, imDir] = imgradient(rgb2gray(im));
5
6 imMag = uint8(imMag);
7
8 figure, imshow(im);
9 figure, imshow(imMag);
```

